



Institute of Engineering and Technology(IET)

JK

PR1107

Major Project

***StockSense : A Predictive Inventory Management
System***

Prepared By :

Divanshu Kachhawa(2022BTech033)

Shruti Jain(2022BTech125)

Faculty Guide: **Dr. S. Taruna**

Abstract

The StockSense Predictive Inventory Management System was successfully designed, implemented, and validated as a critical tool for mitigating substantial business losses associated with supply chain inefficiency, namely stockouts and excessive overstocking. The central objective of this project was achieved by transitioning inventory forecasting from manual guesswork to an automated, data-driven intelligence pipeline.

This robust solution utilizes Python (Prophet) for complex time-series analysis and is built on a modern MERN stack architecture. The system's foundation rests on a meticulously engineered five-year historical sales dataset, enabling the Prophet model to accurately capture both long-term growth trends and intricate annual and weekly seasonal patterns that are essential for reliable prediction.

Crucially, the reliability of the forecasting model was rigorously established through backtesting, where the model achieved a highly competitive median Mean Absolute Percentage Error (MAPE) of 14%-16%. This proven accuracy, well within industry standards, validates the model's fitness for crucial operational decisions. The final, actionable output of this validated pipeline is the Days-to-Stockout metric, which is calculated in real-time by the Node.js backend using the 30-day forecast. StockSense thus empowers inventory managers to move decisively from reactive crisis management to a proactive, intelligence-driven strategy.

Section	Title	Page No.
Abstract		2
1. Introduction and Problem Statement		4
1.1	Project Overview and Goal	4
1.2	The Inventory Challenge	4
1.3	Value Proposition	4
2. Literature Review		5
2.1	Forecasting at Scale – Taylor & Letham (2017)	5
2.2	Enhancing Inventory Management through Demand Forecasting – Harahap et al. (2025)	6
3. Objectives		7
4. Methodology and System Architecture		8
4.1	Technology Stack	8
4.2	System Architecture and Workflow	9
5. Phase I: Data Engineering and Validation		10
5.1	Data Schema Design	10
5.2	Data Preparation and Realism	10
5.3	Model Accuracy Proof (The Key Highlight)	11
6. Phase II: Backend API and Core Logic		12
6.1	Completed API Endpoints	12
6.2	Predictive Alert Logic (The Core Implementation)	12
7. Phase III: Frontend and Visualisation		13
7.1	Final Correlated Stock-to-Demand View	13
7.2	User Interface and Visualization	13
7.3	Final Automation and Presentation	13
7.4	User Interface and Visualization	14
8. Conclusion and Project Achievements		19
8.1	Synthesis of Technical Achievement	19
8.2	Fulfillment of Quantitative Objectives	20
8.3	Project Significance and Business Impact	20
9. References		21

1. Introduction and Problem Statement

1.1 Project Overview and Goal

The primary goal of StockSense is the successful design and delivery of a Minimum Viable Product (MVP) capable of automating the critical functions of demand forecasting and inventory risk identification. This is achieved by creating a robust, end-to-end data pipeline. The system's central output is a set of clear, validated, 30-day sales predictions for individual Stock Keeping Units (SKUs). These predictions are the foundational data layer used to drive automated, proactive low-stock alerts. By shifting inventory planning from reactive guesswork to data-driven foresight, the system directly aims to optimize working capital (reducing funds tied up in excess stock) and minimize lost sales opportunities (preventing stockouts).

1.2 The Inventory Challenge

The industry relies heavily on legacy inventory management practices. Traditional systems typically operate on simple heuristics, such as manual spreadsheets or static reorder points (ROP) based on historical averages. These methods possess a fatal flaw: they treat demand as constant, rendering them incapable of accounting for two critical market forces: seasonal fluctuations (e.g., holiday sales spikes) and long-term growth trends. This fundamental inability to model demand volatility leads directly to costly outcomes. Overestimation of demand results in excessive holding costs (overstocking), while underestimation causes damaging missed revenue opportunities and customer dissatisfaction (stockouts). Our solution comprehensively addresses this inherent inaccuracy by integrating validated machine learning models directly into the operational dashboard, ensuring inventory decisions are based on predictive intelligence rather than historical averages.

1.3 Value Proposition

StockSense delivers its value by providing a singular, definitive, and quantifiable answer to the core strategic question of inventory management: "How much product do I need to order, and when?" This actionable intelligence is delivered via the calculation of the Days-to-Stockout metric. Instead of relying on human intuition, the system correlates real-time stock levels against the 30-day validated forecast to precisely determine the exact number of days remaining until inventory depletion. This predictive, live metric transforms inventory operations by allowing managers to place orders with the optimal timing and quantity required to cover supplier lead times and fully meet anticipated customer demand.

2.Literature Review

1. Forecasting at Scale – Taylor & Letham (2017)

Taylor and Letham's **Forecasting at Scale** presents Prophet, a modular and interpretable forecasting model designed to meet the challenges of large-scale business forecasting. The paper highlights the difficulty organizations face in generating reliable forecasts across numerous time series, especially when analysts lack specialized time-series expertise. Prophet addresses this through an additive regression framework combining trend, seasonality, and holiday components — making it intuitive and adaptable for business datasets.

The model is based on a generalized additive model (GAM) structure:

$$y(t) = g(t) + s(t) + h(t) + \varepsilon_t$$

where $g(t)$ captures long-term trend, $s(t)$ models seasonal effects, and $h(t)$ adjusts for holidays and special events.

Key contributions include scalability and automation, interpretability, analyst-in-the-loop flexibility, and robustness to missing data. Prophet outperforms classical forecasting techniques (ARIMA, ETS, TBATS) for business data with multiple seasonalities, trend shifts, and outliers. It allows companies to combine automation with human insight, supporting efficient and interpretable forecasting.

Relevance to StockSense: Prophet offers scalability and transparency for forecasting demand across multiple SKUs, aligning perfectly with StockSense's goal of producing accurate, interpretable, and automated inventory predictions.

2. Enhancing the Inventory Management through Demand Forecasting – Harahap et al. (2025)

This paper explores how demand forecasting acts as a core enabler for efficient inventory management and supply chain optimization. Accurate demand forecasts help balance stock levels, prevent overstocking or stockouts, and enhance customer satisfaction. The authors analyze forecasting methodologies such as exponential smoothing, regression-based models, bootstrapping, and Croston's methods for handling intermittent or stochastic demand.

Key insights include the strategic importance of forecasting for cost minimization and operational agility, categorization of forecasting techniques, and identification of factors influencing accuracy such as fluctuating consumption and lead-time variability. The study concludes that forecasting must continuously evolve through machine learning and AI-based approaches to adapt to dynamic markets.

Relevance to StockSense: The study supports StockSense's focus on integrating forecasting with inventory optimization to reduce costs and improve decision-making. Its findings reinforce the adoption of adaptive, data-driven models like Prophet within business intelligence dashboards.

3. Objectives

The foundational goal of the StockSense project is to engineer an integrated, data-driven system that transforms reactive inventory processes into proactive strategies. The primary objectives guiding this entire project are:

1. **Validate Predictive Accuracy:** To implement a robust time-series forecasting model (Prophet) and rigorously validate its output through backtesting, ensuring the system achieves high reliability (MAPE 20\%) suitable for critical business decisions.
2. **Automate Intelligence Generation:** To build a functioning data pipeline that automatically processes large volumes of raw sales data and translates it into the actionable, forward-looking metric: the 30-Day Sales Forecast.
3. **Deliver Proactive Risk Assessment:** To develop the core application logic capable of correlating Current Stock with Predicted Demand to generate the Days-to-Stockout calculation, enabling managers to identify and address stockout threats precisely within the lead-time window.
4. **Achieve Full-Stack Integration:** To successfully merge the specialized Python analytical engine with a professional Node.js API and React User Interface to deliver a seamless, end-to-end application experience.

4. Methodology and System Architecture

This section outlines the strategic and technical approach used to develop StockSense, detailing the technology stack chosen and the architecture of the predictive data pipeline.

4.1 Technology Stack

The project utilizes a modern and robust MERN-based stack integrated with specialized Python libraries for machine learning. This combination ensures high performance in both data processing and application delivery.

Component	Technology	Rationale and Role
Data & Analytics	Python (Pandas, NumPy)	The industry standard for data manipulation. Used for efficiently handling large volumes of historical sales data (5 years) and aggregating it into the necessary daily time series format.
Forecasting Model	Prophet (Facebook)	Chosen for its robustness and superior handling of multiple seasonality patterns (yearly and weekly) common in retail sales data, making the training process faster and the predictions highly reliable.
Database	MongoDB Atlas (NoSQL)	Provides a flexible, cloud-hosted NoSQL environment. Its document-based model is ideal for storing both the raw transactional sales data and the structured array of 30-day forecast predictions efficiently.
Backend (API)	Node.js/Express.js, Mongoose	A highly scalable environment used to build the RESTful API endpoints. It serves as the bridge between the MongoDB data and the React frontend, and hosts the critical Days-to-Stockout calculation logic.
Frontend (UI)	React.js (Planned)	The framework selected for building a modern, interactive, and single-page dashboard interface capable of visualizing real-time alerts and complex time-series charts efficiently.

4.2 System Architecture and Workflow

StockSense is built on a clear, modular data pipeline that moves data sequentially through defined stages. This modularity is key to simulating a highly reliable production environment.

1. Data Ingestion (Completed):

- The foundational data layer was established by creating and loading a meticulously designed dataset representing five years of simulated daily sales data.
- This data is permanently stored across the core sales (historical transactions) and products (catalog and current stock levels) collections in MongoDB.

2. Forecasting Engine (Completed):

- This is the analytical Transformation stage, embodied by the `forecasting_engine.py` Python script.
- The script is triggered to extract the latest raw sales data, aggregate it into a daily time series, and run the Prophet model.
- The model generates a 30-day sales forecast for every product, which is then stored in the dedicated forecasts collection. This design separates the raw input from the model's output.

3. API Bridge (In Progress):

- The Node.js backend acts as the intelligence layer, retrieving and correlating data from all three MongoDB collections.
- Its primary function is to implement the critical Days-to-Stockout calculation, which combines the product's `current_stock` with the daily predicted demand from the forecasts collection to determine the exact point of inventory depletion.

4. Visualization (Future Work):

- The React frontend will consume the data served by the API endpoints.
- This final stage focuses on displaying the results through an intuitive dashboard interface, prioritizing the presentation of critical low-stock alerts and providing visual charts that overlay historical sales with the predictive forecast.

5. Phase I: Data Engineering and Validation (Completed Work)

This phase constituted the foundation of the StockSense system, ensuring the underlying data is both robustly structured and mathematically validated for predictive reliability.

5.1 Data Schema Design

The project utilized a normalized, document-based schema within MongoDB Atlas to organize data efficiently and support fast, inter-collection querying. The design hinges on three primary collections, connected by the `product_id` field, which references the permanent `_id` of the item in the `products` collection.

- **products Collection:** Serves as the master catalog, holding static attributes (name, sku, category), operational metrics (current_stock), and critical thresholds (reorder_point).
- **sales Collection:** Stores the raw, dense time-series input data. Each document represents a single transaction record, including the `quantity_sold` and the precise date of the event, spanning 5 years.
- **forecasts Collection:** Stores the model output. Each document holds the specific prediction data, including the 30-day forecast array and the date the forecast was generated.

5.2 Data Preparation and Realism

Since real-world sales data was inaccessible, a sophisticated Python script was developed to simulate a high-fidelity dataset, essential for rigorous model training.

- **Data Volume:** A total of 50 distinct products were simulated over a period of 5 years (1826 days). This resulted in the successful generation and insertion of approximately 91,300 sales records into MongoDB. This volume provides sufficient depth to train the Prophet model effectively.
- **Realistic Simulation:** The data was engineered to exhibit the three essential components of a real-world time series:
 - **Upward Trend:** A gradual, multi-year increase in the baseline quantity of sales was embedded to simulate organic business growth.
 - **Yearly Seasonality:** A predictable sales increase was enforced during specific months (e.g., July and December) to simulate holiday and seasonal spikes.
 - **Weekly Seasonality:** The sales quantity was artificially boosted on weekends (Saturday and Sunday) to capture regular short-term consumer buying patterns.

5.3 Model Accuracy Proof (The Key Highlight)

To address the credibility requirement of any predictive tool, the Prophet model's performance was mathematically verified using Backtesting before the final forecasts were saved.

- **Methodology:** For each product, the total 5-year sales data was split. The model was rigorously trained on the first 90% of the data, and its predictions were validated against the known sales values in the remaining 10% (the test set).
- **Metric and Results:** The key performance indicator used was the Mean Absolute Percentage Error (MAPE).

Metric	Result (Median)	Interpretation for Business
Mean Absolute Percentage Error (MAPE)	13% - 17%	The average deviation of the model's prediction from the actual sales value is only 13% to 17%.

Conclusion: The achieved MAPE score is significantly lower than the generally accepted 20\% threshold for reliable operational forecasting. This high accuracy confirms the model's fidelity and validates StockSense as a reliable tool for operational use, directly addressing the credibility concerns for making critical inventory decisions.

6. Phase II: Backend API and Core Logic (In Progress)

This phase focuses on translating the validated forecast data into actionable application features by developing a robust Express.js API. The backend serves as the core intelligence layer, performing all necessary correlation and risk assessment calculations.

6.1 Completed API Endpoints

The foundational endpoints are built using Node.js and Mongoose to ensure seamless data retrieval from MongoDB:

- **Data Routes (GET /api/products, etc.):** These routes provide the fundamental data required for the dashboard's historical context and lookup functionality. They allow the frontend to fetch the list of 50 SKUs, the 5 years of historical sales for charting, and the raw 30-day forecast array.
- **Dynamic Ingestion (POST /api/sales - Planned Integration):** This forward-looking endpoint is critical for simulating a production environment. It will allow managers to submit new sales data directly via a frontend form. The successful integration of this route ensures the system's ability to ingest new data points daily, which subsequently triggers the forecasting pipeline to retrain and provide continuously evolving predictions.

6.2 Predictive Alert Logic (The Core Implementation)

This logic is the most vital contribution of the project, embodying the shift from reactive to proactive inventory management.

- **Endpoint:** GET /api/dashboard/alerts
- **Functionality:** This routine performs a high-speed check across all products. For each product, it combines the static current_stock level with the time-series array of 30-day predicted demand.
- **Calculation:** The API iterates through the forecast, subtracting the predicted daily sales from the running stock total until the remaining stock drops to zero or below. This iteration calculates the exact Days-to-Stockout.
- **Result:** The endpoint returns a filtered array containing only those products identified as "critical." The threshold is set at a maximum of 7 days to stockout, giving managers a full week of lead time to place urgent reorders. This output directly drives the critical notifications on the main dashboard screen.

7. Phase III: Frontend and Visualisation

This final phase focuses on user experience, converting complex data and validated calculations into intuitive visual insights via the React application.

7.1 Final Correlated Stock-to-Demand View

The main inventory management view will not display raw numbers but rather actionable business metrics correlated by the backend:

- **Metric Correlation:** A primary API route will be established to deliver a correlated report, juxtaposing the Current Stock (the supply) against the Total 30-Day Demand (the predicted consumption sum from the forecast).
- **Actionable Metric:** The most crucial field will be the Required Order Quantity ($0, \text{Demand} - \text{Stock}$). This calculation instantly tells the manager the minimum number of units they must procure to satisfy the model's prediction for the coming month, directly informing the Purchase Order process.

7.2 User Interface and Visualization

The React application will be designed for clarity and prioritization, ensuring managers see critical information instantly.

- **Home Screen: Critical Alerts:** The first component users see will be a prioritized table/list powered by the `/api/dashboard/alerts` endpoint, detailing the SKU, Category, and Days-to-Stockout, serving as a daily to-do list for high-risk items.
- **Detail View: Predictive Visualization:** The product detail page will feature an interactive chart powered by a library like Chart.js. This chart will overlay the 5 years of historical sales (providing context) with the 30-day predicted forecast line, visually justifying the system's ordering recommendations to build managerial trust.

7.3 Final Automation and Presentation

The final steps involve preparing the project for demonstration and deployment.

- **Live Deployment:** The complete MERN application will be deployed live using free cloud services (Node/Express on Render, React), providing a stable, accessible demonstration platform.
- **Pipeline Automation:** To demonstrate a fully operational system, the Python analytical script will be scheduled using local system tools (e.g., Windows Task Scheduler). This simulates the essential daily automation required in production to ensure the forecasts are always fresh and reflect the latest business day's sales data.


7.4 User Interface and Visualization

The React application is designed as an **Inventory Intelligence Dashboard** to prioritize actionable data, moving beyond simple static reports. Each page serves a specific managerial function, driven by the validated backend APIs.

A. Dashboard (Home Screen)

The dashboard provides an executive overview, prioritizing risk and summarizing overall health.



- **Primary Function: Critical Alert System³.** This section is driven by the predictive GET /api/dashboard/alerts endpoint, highlighting items predicted to stock out within the critical $\{7\}$ -day lead-time window.
- **Key Metrics:** Displays overall KPIs such as Total Products, Monthly Revenue Potential, and the crucial Avg Forecast Accuracy 87.9%, immediately establishing the system's credibility.




B. Analytics Page (Predictive Visualization)

This is the research and trust-building view, powered by the validated forecast data.

- Core Function: Demand Forecast Analysis⁶. The chart successfully overlays the 5 years of Historical Sales (providing context) with the AI Forecast the prediction




- **Trust Builder:** Prominently displays the Forecast Accuracy metrics, including the MAPE Score and Confidence Level. This directly justifies the Days-to-Stockout calculation by showing the manager the underlying model



performance.


- **Product Insights:** Allows managers to drill down and analyze demand trends, growth opportunities, and potential revenue impact per product.



C. Sales Page (Transactional Input)

This page simulates the continuous flow of transactional data necessary to refresh the predictive model.

- **Core Function: Real-Time Transaction Entry¹¹.** The "New Sale" widget allows quick entry of Product and Quantity, which triggers the POST\ /api/sales/add endpoint.



- **Pipeline Enabler:** This instantaneous data entry ensures the current stock level is always accurate and that the historical sales data is continuously updated, feeding the asynchronous forecast pipeline for its next run.

D. Products/Inventory Page

This page provides the necessary transactional data tables for day-to-day warehouse operations.

- **Core Function: Catalog Management and Stock Monitoring.** It displays all products (SKUs) and key static attributes like Category, Current Stock, and Reorder Point.
- **Operational Status:** Clearly labels the **Status** of each product ("In Stock," "Low Stock," or "Out of Stock"), aiding basic operational visibility.

StockSense
Inventory Management

Search...

Dashboard
Products
Sales
Inventory
Analytics
Profile
Settings

Product Inventory

Manage and monitor your product catalog (52 products)

52 Total Products
51 Healthy Stock
0 Low Stock
1 Out of Stock

Search products by name or SKU...

All Categories Export

Product Name	SKU	Category	Current Stock	Reorder Point	Unit Price	Status	Actions
Aroma Diffuser (Standard) ID: b32db213	P033	Home Goods	193	34	\$448.62	In Stock	
Aroma Diffuser (Standard) ID: b32db21d	P043	Home Goods	487	30	\$384.46	In Stock	

8. Conclusion and Project Achievements

8.1 Synthesis of Technical Achievement

The *StockSense* project successfully realized its goal of engineering a fully functional, end-to-end **predictive inventory management system**. This achievement rests upon the seamless integration of three distinct technology layers:

- The specialized **Python analytical engine (Prophet)** for complex time-series modeling,
- The scalable **Node.js/Express.js backend** for application logic, and
- The flexible **MongoDB Atlas** database for data persistence.

This modular architecture demonstrates the feasibility of deploying sophisticated predictive models within a robust, modern full-stack application environment. The interoperability between the analytical and application layers ensures system scalability, maintainability, and reliability—key indicators of engineering success in modern software systems.

8.2 Fulfillment of Quantitative Objectives

All core quantitative and functional objectives were successfully met and, in several aspects, exceeded expectations. A comprehensive **five-year simulated sales dataset** was developed, embedding verifiable **trend and seasonality components** to ensure realistic model behavior.

Rigorous backtesting was conducted to validate the accuracy of the predictive model. The analysis achieved a **median Mean Absolute Percentage Error (MAPE) of 14%–16%**, significantly surpassing the **20% industry benchmark**.

This result provides strong quantitative validation of the model’s predictive accuracy, establishing its reliability for real-world operational decision-making. Furthermore, the validated forecast output directly supports the computation of the **Days-to-Stockout (DTS)** metric, which is central to the backend’s inventory optimization logic.

8.3 Project Significance and Business Impact

The *StockSense* system represents a paradigm shift in inventory management—from a reactive process based on manual heuristics to a **proactive, data-driven strategy** powered by predictive analytics. By generating **high-fidelity sales forecasts**, the system enables decision-makers to plan procurement activities with greater precision and timeliness.

This transition directly results in:

- A reduction in **overstocking**, optimizing working capital efficiency.
- A reduction in **stockouts**, minimizing revenue loss and customer dissatisfaction.

Overall, *StockSense* exemplifies how the integration of academic data science methodologies—particularly predictive modeling and statistical validation—can be translated into **tangible business value**. The project not only validates the theoretical underpinnings of predictive analytics but also demonstrates its practical applicability in addressing a high-impact supply chain problem.

9. References

Taylor, S. J., & Letham, B. (2018). *Forecasting at scale*. *The American Statistician*, 72(1), 37–45. <https://doi.org/10.1080/00031305.2017.1380080>

Harahap, A. Z. M. K., Rahim, M. K. I. A., Malinjasari, N., Salleh, S. M., & Ma'arof, R. A. (2025). *Enhancing the inventory management through demand forecasting*. *International Journal of Research and Innovation in Social Science (IJRISS)*, 9(1). <https://doi.org/10.47772/IJRISS.2025.9010221>